

MANAV RACHNA UNIVERSITY
SCHOOL OF ENGINEERING
DEPARTMENT OF COMPUTER SCIENCE & TECHNOLOGY

LAB FILE

Supervised Learning (CSH212B-P/ CSH212B-T)

Submitted to:

Dr. ROSHI SAXENA

Manager - Xebia Academy

Submitted by:

Devanaboina Tharun

2K23CSUN01247 AIML-
3B

MANAV RACHNA UNIVERSITY

SCHOOL OF ENGINEERING

DEPARTMENT OF COMPUTER SCIENCE & TECHNOLOGY

Supervised Learning Projects: - 3B

S. No	Name of the Program	Date
1	Write a python code to demonstrate commands for numpy and pandas.	
2	Write a python program to calculate mean square and mean absolute error.	
3	Write a python program to calculate gradient descent of a machine learning model.	
4	Prepare a linear regression model for predicting the salary of user based on number of years of experience.	
5	Prepare a linear regression model for prediction of resale car price.	
6	Prepare a Lasso and Ridge regression model for prediction of house price and compare it with linear regression model.	
7	Prepare a decision tree model for Iris Dataset using Gini Index.	
8	Prepare a decision tree model for Iris Dataset using entropy.	
9	Prepare a naïve bayes classification model for prediction of purchase power of a user.	
10	Prepare a naïve bayes classification model for classification of email messages into spam or not spam.	
11	Prepare a model for prediction of prostate cancer using KNN Classifier.	
12	Prepare a model for prediction of survival from Titanic Ship using Random Forest and compare the accuracy with other classifiers also.	

□ LAB-01

#question-1)Write a python code to demonstrate commands for numpy and pandas

```
import numpy as np
import pandas as pd
```

NumPy


```
arr = np.array([1, 2, 3, 4, 5])
matrix = np.array([[1, 2, 3], [4, 5, 6]])
mean = np.mean(arr)
std_dev = np.std(arr)
```

```
print("NumPy Array:", arr) print("Matrix:\n",
matrix) print("Mean of Array:", mean)
print("Standard Deviation of Array:", std_dev)
```

Pandas

```
data = pd.Series([10, 20, 30, 40], index=['a', 'b', 'c', 'd'])
df = pd.DataFrame({'Name': ['Alice', 'Bob', 'Charlie'], 'Age': [25,
df['Pass'] = df['Score'] > 50
```

```
print("\nPandas Series:\n", data)
print("\nPandas DataFrame:\n", df)
```

 NumPy Array: [1 2 3 4 5] Matrix: [[1 2 3] [4 5 6]]
Mean of Array: 3.0
Standard Deviation of Array: 1.4142135623730951

Pandas Series:

```
a    10
b    20
c    30
d    40
dtype: int64
```

Pandas DataFrame:

	Name	Age	Score	Pass
0	Alice	25	85	True
1	Bob	30	90	True
2	Charlie	35	95	True

□ LAB-02

#question-2)Write a python program to calculate mean square and mean absolute error

```
def calculate_mse_and_mae(actual, predicted):
    n = len(actual)
```

```
    sum_squared_error = 0
    sum_absolute_error = 0
```

```

for i in range(n):
    sum_squared_error += (actual[i] - predicted[i]) ** 2
    sum_absolute_error += abs(actual[i] - predicted[i])

mse = sum_squared_error / n
mae = sum_absolute_error / n

return mse, mae

```

```

actual = [3, -0.5, 2, 7]
predicted = [2.5, 0.0, 2, 8]

```

```

mse, mae = calculate_mse_and_mae(actual, predicted)
print(f"Mean Squared Error: {mse}")
print(f"Mean Absolute Error: {mae}")

```

LAB-03

#question-3)Write a python program to calculate gradient descent of

```

import numpy as np
import matplotlib.pyplot as plt

```

```

def plot_function(func):
    x_values = np.linspace(-10, 10, 1000)
    y_values = [func(x) for x in x_values]
    print("x values are: ", x_values)
    print("y values are: ", y_values)
    plt.plot(x_values, y_values)
    plt.show()
plot_function(lambda x: np.sin(x))


```

↩ x values are: [-10. -9.97997998 -9.95995996 -9.93993994 -9.91991992
-9.8998999 -9.87987988 -9.85985986 -9.83983984 -9.81981982
-9.7997998 -9.77977978 -9.75975976 -9.73973974 -9.71971972
-9.6996997 -9.67967968 -9.65965966 -9.63963964 -9.61961962
-9.5995996 -9.57957958 -9.55955956 -9.53953954 -9.51951952
-9.4994995 -9.47947948 -9.45945946 -9.43943944 -9.41941942
-9.3993994 -9.37937938 -9.35935936 -9.33933934 -9.31931932
-9.2992993 -9.27927928 -9.25925926 -9.23923924 -9.21921922
-9.1991992 -9.17917918 -9.15915916 -9.13913914 -9.11911912
-9.0990991 -9.07907908 -9.05905906 -9.03903904 -9.01901902
-8.998999 -8.97897898 -8.95895896 -8.93893894 -8.91891892
-8.8988989 -8.87887888 -8.85885886 -8.83883884 -8.81881882
-8.7987988 -8.77877878 -8.75875876 -8.73873874 -8.71871872
-8.6986987 -8.67867868 -8.65865866 -8.63863864 -8.61861862
-8.5985986 -8.57857858 -8.55855856 -8.53853854 -8.51851852
-8.4984985 -8.47847848 -8.45845846 -8.43843844 -8.41841842
-8.3983984 -8.37837838 -8.35835836 -8.33833834 -8.31831832
-8.2982983 -8.27827828 -8.25825826 -8.23823824 -8.21821822
-8.1981982 -8.17817818 -8.15815816 -8.13813814 -8.11811812
-8.0980981 -8.07807808 -8.05805806 -8.03803804 -8.01801802
-7.997998 -7.97797798 -7.95795796 -7.93793794 -7.91791792
-7.8978979 -7.87787788 -7.85785786 -7.83783784 -7.81781782
-7.7977978 -7.77777778 -7.75775776 -7.73773774 -7.71771772
-7.6976977 -7.67767768 -7.65765766 -7.63763764 -7.61761762
-7.5975976 -7.57757758 -7.55755756 -7.53753754 -7.51751752
-7.4974975 -7.47747748 -7.45745746 -7.43743744 -7.41741742
-7.3973974 -7.37737738 -7.35735736 -7.33733734 -7.31731732
-7.2972973 -7.27727728 -7.25725726 -7.23723724 -7.21721722
-7.1971972 -7.17717718 -7.15715716 -7.13713714 -7.11711712
-7.0970971 -7.07707708 -7.05705706 -7.03703704 -7.01701702
-6.996997 -6.97697698 -6.95695696 -6.93693694 -6.91691692
-6.8968969 -6.87687688 -6.85685686 -6.83683684 -6.81681682
-6.7967968 -6.77677678 -6.75675676 -6.73673674 -6.71671672
-6.6966967 -6.67667668 -6.65665666 -6.63663664 -6.61661662
-6.5965966 -6.57657658 -6.55655656 -6.53653654 -6.51651652
-6.4964965 -6.47647648 -6.45645646 -6.43643644 -6.41641642
-6.3963964 -6.37637638 -6.35635636 -6.33633634 -6.31631632
-6.2962963 -6.27627628 -6.25625626 -6.23623624 -6.21621622
-6.1961962 -6.17617618 -6.15615616 -6.13613614 -6.11611612
-6.0960961 -6.07607608 -6.05605606 -6.03603604 -6.01601602
-5.995996 -5.97597598 -5.95595596 -5.93593594 -5.91591592
-5.8958959 -5.87587588 -5.85585586 -5.83583584 -5.81581582
-5.7957958 -5.77577578 -5.75575576 -5.73573574 -5.71571572
-5.6956957 -5.67567568 -5.65565566 -5.63563564 -5.61561562
-5.5955956 -5.57557558 -5.55555556 -5.53553554 -5.51551552
-5.4954955 -5.47547548 -5.45545546 -5.43543544 -5.41541542
-5.3953954 -5.37537538 -5.35535536 -5.33533534 -5.31531532
-5.2952953 -5.27527528 -5.25525526 -5.23523524 -5.21521522
-5.1951952 -5.17517518 -5.15515516 -5.13513514 -5.11511512
-5.0950951 -5.07507508 -5.05505506 -5.03503504 -5.01501502
-4.99499499 -4.97497497 -4.95495495 -4.93493493 -4.91491491
-4.89489489 -4.87487487 -4.85485485 -4.83483483 -4.81481481
-4.79479479 -4.77477477 -4.75475475 -4.73473473 -4.71471471
-4.69469469 -4.67467467 -4.65465465 -4.63463463 -4.61461461
-4.59459459 -4.57457457 -4.55455455 -4.53453453 -4.51451451
-4.49449449 -4.47447447 -4.45445445 -4.43443443 -4.41441441
-4.39439439 -4.37437437 -4.35435435 -4.33433433 -4.31431431
-4.29429429 -4.27427427 -4.25425425 -4.23423423 -4.21421421
-4.19419419 -4.17417417 -4.15415415 -4.13413413 -4.11411411
-4.09409409 -4.07407407 -4.05405405 -4.03403403 -4.01401401
-3.99399399 -3.97397397 -3.95395395 -3.93393393 -3.91391391
-3.89389389 -3.87387387 -3.85385385 -3.83383383 -3.81381381
-3.79379379 -3.77377377 -3.75375375 -3.73373373 -3.71371371
-3.69369369 -3.67367367 -3.65365365 -3.63363363 -3.61361361
-3.59359359 -3.57357357 -3.55355355 -3.53353353 -3.51351351
-3.49349349 -3.47347347 -3.45345345 -3.43343343 -3.41341341
-3.39339339 -3.37337337 -3.35335335 -3.33333333 -3.31331331
-3.29329329 -3.27327327 -3.25325325 -3.23323323 -3.21321321
-3.19319319 -3.17317317 -3.15315315 -3.13313313 -3.11311311
-3.09309309 -3.07307307 -3.05305305 -3.03303303 -3.01301301
-2.99299299 -2.97297297 -2.95295295 -2.93293293 -2.91291291
-2.89289289 -2.87287287 -2.85285285 -2.83283283 -2.81281281
-2.79279279 -2.77277277 -2.75275275 -2.73273273 -2.71271271
-2.69269269 -2.67267267 -2.65265265 -2.63263263 -2.61261261
-2.59259259 -2.57257257 -2.55255255 -2.53253253 -2.51251251
-2.49249249 -2.47247247 -2.45245245 -2.43243243 -2.41241241
-2.39239239 -2.37237237 -2.35235235 -2.33233233 -2.31231231
-2.29229229 -2.27227227 -2.25225225 -2.23223223 -2.21221221
-2.19219219 -2.17217217 -2.15215215 -2.13213213 -2.11211211
-2.09209209 -2.07207207 -2.05205205 -2.03203203 -2.01201201
-1.99199199 -1.97197197 -1.95195195 -1.93193193 -1.91191191
-1.89189189 -1.87187187 -1.85185185 -1.83183183 -1.81181181
-1.79179179 -1.77177177 -1.75175175 -1.73173173 -1.71171171
-1.69169169 -1.67167167 -1.65165165 -1.63163163 -1.61161161
-1.59159159 -1.57157157 -1.55155155 -1.53153153 -1.51151151
-1.49149149 -1.47147147 -1.45145145 -1.43143143 -1.41141141
-1.39139139 -1.37137137 -1.35135135 -1.33133133 -1.31131131
-1.29129129 -1.27127127 -1.25125125 -1.23123123 -1.21121121
-1.19119119 -1.17117117 -1.15115115 -1.13113113 -1.11111111
-1.09109109 -1.07107107 -1.05105105 -1.03103103 -1.01101101

-0.99099099	-0.97097097	-0.95095095	-0.93093093	-0.91091091
-0.89089089	-0.87087087	-0.85085085	-0.83083083	-0.81081081
-0.79079079	-0.77077077	-0.75075075	-0.73073073	-0.71071071
-0.69069069	-0.67067067	-0.65065065	-0.63063063	-0.61061061
-0.59059059	-0.57057057	-0.55055055	-0.53053053	-0.51051051
-0.49049049	-0.47047047	-0.45045045	-0.43043043	-0.41041041
-0.39039039	-0.37037037	-0.35035035	-0.33033033	-0.31031031
-0.29029029	-0.27027027	-0.25025025	-0.23023023	-0.21021021
-0.19019019	-0.17017017	-0.15015015	-0.13013013	-0.11011011
-0.09009009	-0.07007007	-0.05005005	-0.03003003	-0.01001001
0.01001001	0.03003003	0.05005005	0.07007007	0.09009009
0.11011011	0.13013013	0.15015015	0.17017017	0.19019019
0.21021021	0.23023023	0.25025025	0.27027027	0.29029029
0.31031031	0.33033033	0.35035035	0.37037037	0.39039039
0.41041041	0.43043043	0.45045045	0.47047047	0.49049049
0.51051051	0.53053053	0.55055055	0.57057057	0.59059059
0.61061061	0.63063063	0.65065065	0.67067067	0.69069069
0.71071071	0.73073073	0.75075075	0.77077077	0.79079079
0.81081081	0.83083083	0.85085085	0.87087087	0.89089089
0.91091091	0.93093093	0.95095095	0.97097097	0.99099099
1.01101101	1.03103103	1.05105105	1.07107107	1.09109109
1.11111111	1.13113113	1.15115115	1.17117117	1.19119119
1.21121121	1.23123123	1.25125125	1.27127127	1.29129129
1.31131131	1.33133133	1.35135135	1.37137137	1.39139139
1.41141141	1.43143143	1.45145145	1.47147147	1.49149149
1.51151151	1.53153153	1.55155155	1.57157157	1.59159159
1.61161161	1.63163163	1.65165165	1.67167167	1.69169169
1.71171171	1.73173173	1.75175175	1.77177177	1.79179179
1.81181181	1.83183183	1.85185185	1.87187187	1.89189189
1.91191191	1.93193193	1.95195195	1.97197197	1.99199199
2.01201201	2.03203203	2.05205205	2.07207207	2.09209209
2.11211211	2.13213213	2.15215215	2.17217217	2.19219219
2.21221221	2.23223223	2.25225225	2.27227227	2.29229229
2.31231231	2.33233233	2.35235235	2.37237237	2.39239239
2.41241241	2.43243243	2.45245245	2.47247247	2.49249249
2.51251251	2.53253253	2.55255255	2.57257257	2.59259259
2.61261261	2.63263263	2.65265265	2.67267267	2.69269269
2.71271271	2.73273273	2.75275275	2.77277277	2.79279279
2.81281281	2.83283283	2.85285285	2.87287287	2.89289289
2.91291291	2.93293293	2.95295295	2.97297297	2.99299299
3.01301301	3.03303303	3.05305305	3.07307307	3.09309309
3.11311311	3.13313313	3.15315315	3.17317317	3.19319319
3.21321321	3.23323323	3.25325325	3.27327327	3.29329329
3.31331331	3.33333333	3.35335335	3.37337337	3.39339339
3.41341341	3.43343343	3.45345345	3.47347347	3.49349349
3.51351351	3.53353353	3.55355355	3.57357357	3.59359359
3.61361361	3.63363363	3.65365365	3.67367367	3.69369369
3.71371371	3.73373373	3.75375375	3.77377377	3.79379379
3.81381381	3.83383383	3.85385385	3.87387387	3.89389389
3.91391391	3.93393393	3.95395395	3.97397397	3.99399399
4.01401401	4.03403403	4.05405405	4.07407407	4.09409409
4.11411411	4.13413413	4.15415415	4.17417417	4.19419419
4.21421421	4.23423423	4.25425425	4.27427427	4.29429429
4.31431431	4.33433433	4.35435435	4.37437437	4.39439439
4.41441441	4.43443443	4.45445445	4.47447447	4.49449449
4.51451451	4.53453453	4.55455455	4.57457457	4.59459459
4.61461461	4.63463463	4.65465465	4.67467467	4.69469469
4.71471471	4.73473473	4.75475475	4.77477477	4.79479479
4.81481481	4.83483483	4.85485485	4.87487487	4.89489489
4.91491491	4.93493493	4.95495495	4.97497497	4.99499499
5.01501502	5.03503504	5.05505506	5.07507508	5.0950951
5.11511512	5.13513514	5.15515516	5.17517518	5.1951952
5.21521522	5.23523524	5.25525526	5.27527528	5.2952953
5.31531532	5.33533534	5.35535536	5.37537538	5.3953954
5.41541542	5.43543544	5.45545546	5.47547548	5.4954955
5.51551552	5.53553554	5.55555556	5.57557558	5.5955956
5.61561562	5.63563564	5.65565566	5.67567568	5.6956957
5.71571572	5.73573574	5.75575576	5.77577578	5.7957958
5.81581582	5.83583584	5.85585586	5.87587588	5.8958959
5.91591592	5.93593594	5.95595596	5.97597598	5.995996
6.01601602	6.03603604	6.05605606	6.07607608	6.0960961
6.11611612	6.13613614	6.15615616	6.17617618	6.1961962
6.21621622	6.23623624	6.25625626	6.27627628	6.2962963
6.31631632	6.33633634	6.35635636	6.37637638	6.3963964
6.41641642	6.43643644	6.45645646	6.47647648	6.4964965
6.51651652	6.53653654	6.55655656	6.57657658	6.5965966
6.61661662	6.63663664	6.65665666	6.67667668	6.6966967
6.71671672	6.73673674	6.75675676	6.77677678	6.7967968
6.81681682	6.83683684	6.85685686	6.87687688	6.8968969
6.91691692	6.93693694	6.95695696	6.97697698	6.996997
7.01701702	7.03703704	7.05705706	7.07707708	7.0970971
7.11711712	7.13713714	7.15715716	7.17717718	7.1971972
7.21721722	7.23723724	7.25725726	7.27727728	7.2972973
7.31731732	7.33733734	7.35735736	7.37737738	7.3973974
7.41741742	7.43743744	7.45745746	7.47747748	7.4974975
7.51751752	7.53753754	7.55755756	7.57757758	7.5975976
7.61761762	7.63763764	7.65765766	7.67767768	7.6976977
7.71771772	7.73773774	7.75775776	7.77777778	7.7977978
7.81781782	7.83783784	7.85785786	7.87787788	7.8978979
7.91791792	7.93793794	7.95795796	7.97797798	7.997998

8.01801802	8.03803804	8.05805806	8.07807808	8.0980981
8.11811812	8.13813814	8.15815816	8.17817818	8.1981982
8.21821822	8.23823824	8.25825826	8.27827828	8.2982983
8.31831832	8.33833834	8.35835836	8.37837838	8.3983984
8.41841842	8.43843844	8.45845846	8.47847848	8.4984985
8.51851852	8.53853854	8.55855856	8.57857858	8.5985986
8.61861862	8.63863864	8.65865866	8.67867868	8.6986987
8.71871872	8.73873874	8.75875876	8.77877878	8.7987988
8.81881882	8.83883884	8.85885886	8.87887888	8.8988989
8.91891892	8.93893894	8.95895896	8.97897898	8.998999
9.01901902	9.03903904	9.05905906	9.07907908	9.0990991
9.11911912	9.13913914	9.15915916	9.17917918	9.1991992
9.21921922	9.23923924	9.25925926	9.27927928	9.2992993
9.31931932	9.33933934	9.35935936	9.37937938	9.3993994
9.41941942	9.43943944	9.45945946	9.47947948	9.4994995
9.51951952	9.53953954	9.55955956	9.57957958	9.5995996
9.61961962	9.63963964	9.65965966	9.67967968	9.6996997
9.71971972	9.73973974	9.75975976	9.77977978	9.7997998
9.81981982	9.83983984	9.85985986	9.87987988	9.8998999
9.91991992	9.93993994	9.95995996	9.97997998	10.

y values are: [0.5440211108893698, 0.5271149856654235, 0.5099975991781289, 0.49267581186741427, 0.4751565660945675, 0.4574468833598]



```
plot_derivative(lambda x: np.sin(x))
```

[illegible]

▲


```

5 2.9 6 3.0 7 3.2
8 3.2
9 3.7
10 3.9
11 4.0 12 4.0 13 4.1
14 4.5
15 4.9
16 5.1
17 5.3 18 5.9 19 6.0
20 6.8
21 7.1
22 7.9
23 8.2
24 8.7
25 9.0
26 9.5
27 9.6
28 10.3
29 10.5
Name: YearsExperience, dtype: float64)

```

```

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3, ra
x_train,x_test,y_train,y_test

```

```

6 60150.0
18 81363.0
13 57081.0
7 54445.0
27 112635.0
1 46205.0
16 66029.0
0 39343.0
15 67938.0
29 121872.0
28 122391.0
9 57189.0
8 64445.0
12 56957.0
11 55794.0
5 56642.0,
Salary
17 83088.0
21 98273.0
10 63218.0
19 93940.0
14 61111.0
20 91738.0
26 116969.0
3 43525.0
24 109431.0,
22 7.9
23 8.2
4 2 2.2
25 1.5
6 9.0
18 3.0
13 5.9
7 4.1
27 3.2
1 9.6
16 1.3
0 5.1
15 1.1
29 4.9
28 10.5
9 10.3
8 3.7
12 3.2
11 4.0
5 4.0
2.9

```

```

Name: YearsExperience, dtype: float64,
17 5.3
21 7.1

```

```
from sklearn.linear_model import LinearRegression  
model=LinearRegression()  
model.fit(x_train,y_train)
```

LinearRegression

LinearRegression()

```
model.coef_
```

array([0.00010441])

```
model.intercept_
```

-2.522510616511819

```
model.score(x_test,y_test)
```

0.9242662549548135

```
x_pred=model.predict(x_train)  
y_pred=model.predict(x_test)  
x_pred,y_pred
```

(array([8.05410626, 9.36023523, 1.64238074, 1.41686206, 8.50096733,
3.75755794, 5.97233923, 3.4371335 , 3.16191719, 9.23734844,
2.30160522, 4.37136547, 1.58516581, 4.57067804, 10.20175398,
10.2559411 , 3.44840943, 4.20598511, 3.42418705, 3.30276195,
3.39129891]),
array([6.15244094, 7.73785808, 4.07787798, 7.28546345, 3.85789287,
7.0555597 , 9.68984747, 2.02179502, 8.90282907]))

```
plt.scatter(x_train,y_train,color="green")  
plt.plot(x_train,x_pred,color="blue")  
plt.ylabel("Years Of Experience")  
plt.xlabel("Salary")  
plt.title("Salary vs Experience")  
plt.show()
```



LAB-05

#QUESTION-5) Prepare a linear regression model for prediction of resale

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
df3=pd.read_csv("/content/cars24-car-price-cleaned.csv")
```

```
df3.info()
```

```
df3.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19820 entries, 0 to 19819
Data columns (total 18 columns):
 #   Non-Null Count  Dtype
---  -
0   19820 non-null float64
1   19820 non-null float64
2   19820 non-null int64
3   19820 non-null int64
4   19820 non-null float64
5   19820 non-null float64
6   19820 non-null float64
7   19820 non-null float64
8   19820 non-null float64
9   19820 non-null object
10  19820 non-null object
11  19820 non-null int64
12  19820 non-null int64
13  19820 non-null int64
14  19820 non-null int64
15  19820 non-null int64
16  19820 non-null int64
17  19820 non-null int64
18  19820 non-null int64
dtypes: float64(6), int64(10), object(2)
memory usage: 2.7+ MB
```

	selling_price	year	km_driven	mileage	engine	max_power	age	Individual	Trustmark Dealer
count	19820.000000	19820.000000	1.982000e+04	19820.000000	19820.000000	19820.000000	19820.000000	19820.000000	19820.000000
mean	6.585509	2014.561453	5.815856e+04	19.503402	1475.702381	98.122907	8.438547	0.390666	0.009586
std	4.847364	3.196636	5.171563e+04	4.297784	518.571223	44.761727	3.196636	0.487912	0.097442
min	0.300000	1992.000000	1.000000e+02	4.000000	0.000000	5.000000	2.000000	0.000000	0.000000
25%	3.410000	2013.000000	3.100000e+04	16.950000	1197.000000	73.900000	6.000000	0.000000	0.000000
50%	5.200000	2015.000000	5.200000e+04	19.300000	1248.000000	86.800000	8.000000	0.000000	0.000000
75%	7.850000	2017.000000	7.400000e+04	22.320000	1582.000000	112.000000	10.000000	1.000000	0.000000
max	20.902500	2021.000000	3.800000e+06	120.000000	6752.000000	626.000000	31.000000	1.000000	1.000000

```
df3["make"]=df3.groupby("make")["selling_price"].transform("mean")
```

```
df3["model"]=df3.groupby("model")["selling_price"].transform("mean")
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler=MinMaxScaler()
```

```
df4=pd.DataFrame(scaler.fit_transform(df3),columns=df3.columns)
```

```
df4
```

	selling_price	year	km_driven	mileage	engine	max_power	age	make	model	Individual	Trustmark Dealer	Diesel
0	0.043684	0.689655	0.031553	0.135345	0.117891	0.066506	0.310345	0.194048	0.041550	1.0	0.0	0.0
1	0.252397	0.827586	0.005237	0.128448	0.177281	0.123994	0.172414	0.232517	0.218382	1.0	0.0	0.0
2	0.089795	0.620690	0.015764	0.112069	0.177281	0.120773	0.379310	0.232517	0.149143	1.0	0.0	0.0
3	0.095134	0.689655	0.009711	0.145862	0.147808	0.100000	0.310345	0.194048	0.093193	1.0	0.0	0.0
4	0.262104	0.793103	0.007869	0.161810	0.221860	0.150709	0.206897	0.252367	0.313574	0.0	0.0	1.0
...
19815	0.300934	0.862069	0.018258	0.168879	0.202014	0.099919	0.137931	0.484670	0.328028	0.0	0.0	1.0
19816	0.434413	0.931034	0.004711	0.116379	0.203347	0.138647	0.068966	0.194048	0.330632	0.0	0.0	0.0
19817	0.191724	0.793103	0.017606	0.147759	0.221860	0.158647	0.206897	0.318156	0.200656	0.0	0.0	1.0
19818	0.580027	0.827586	1.000000	0.103448	0.322719	0.217391	0.172414	0.324782	0.377671	0.0	0.0	1.0
19819	0.567892	0.931034	0.003395	0.120690	0.221712	0.181320	0.068966	0.258412	0.519465	0.0	0.0	0.0

19820 rows × 13 columns

Next steps:

[Generate code with df4](#)

☐ [View recommended plots](#)

[New interactive sheet](#)

```
from sklearn.model_selection import train_test_split
y=df3["selling_price"]
x=df3.drop("selling_price",axis=1)
y.shape,x.shape
```

((19820,), (19820, 17))

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,ran
x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

((13874, 17), (5946, 17), (13874,), (5946,))

```
from sklearn.linear_model import LinearRegression
model=LinearRegression()
model.fit(x_train,y_train)
```

LinearRegression ⓘ ?

LinearRegression()

```
from sklearn.linear_model import LinearRegression
model= LinearRegression()
model.fit(x_train,y_train)
```

LinearRegression ⓘ ?

LinearRegression()

model.coef_

array([8.94320251e-02, -1.35638241e-06, -4.05906554e-02, 2.29106553e-04,
1.50304468e-03, -8.94320251e-02, 6.61474952e-02, 8.61386888e-01,
-1.44630798e-01, -1.44855016e-01, 1.38520375e-01, 2.66216225e+00,
3.30456609e-01, -1.36368445e-01, -8.04585280e-02, -3.35811569e-01,
-4.86084483e-01])

model.intercept_

-178.08232225367115

```
x_pred=model.predict(x_train)
y_pred=model.predict(x_test)
x_pred
```

```
→ array([ 2.95898276, 20.91025983, 7.69454424, ..., 3.24968466,
          5.44348659, 2.19881995])
```

```
y_test_predict=model.predict(x_test)
y_test_predict
```


```
→ array([ 1.135053 , 4.78976725, 5.96846351, ..., 1.20471131,
          3.08615105, 10.67406941])
```

```
model.score(x_test,y_test)
```

```
→ 0.9458843076992299
```

```
import matplotlib.pyplot as plt
fig=plt.figure()
plt.scatter(y_test_predict,y_test)
plt.legend()
plt.show()
```

```
→ WARNING:matplotlib.legend:No artists with labels found to put in legend. Note that artists whose label start with an underscore are
```




```
import seaborn as sns
imp=pd.DataFrame(list(zip(x.columns,np.abs(model.coef_))),columns=["
sns.barplot(x="features",y="coefficient",data=imp)
plt.xticks(rotation=90)
```

```

([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16],
 [Text(0, 0, 'year'),
  Text(1, 0, 'km_driven'),
  Text(2, 0, 'mileage'),
  Text(3, 0, 'engine'), Text(4,
0, 'max_power'), Text(5, 0,
'age'), Text(6, 0, 'make'),
Text(7, 0, 'model'),
Text(8, 0, 'Individual'),
Text(9, 0, 'Trustmark Dealer'),
Text(10, 0, 'Diesel'), Text(11,
0, 'Electric'), Text(12, 0,
'LPG'),
Text(13, 0, 'Petrol'),
Text(14, 0, 'Manual'), Text(15,
0, '5'), Text(16, 0, '>5')])

```



LAB-06

```

#QUESTION-6)prepare a l1 and l2 regression models for prediction of
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.linear_model import LinearRegression,Lasso,Ridge
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error

```

```

df=pd.read_csv("housing.csv")
df.head()

```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0

```
df.shape
```

```
(15210, 10)
```

```
df["ocean_proximity"]=df.groupby("ocean_proximity")["median_house_value"].transform("median")
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0	8.3252	452600.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0	8.3014	358500.0
2	-122.25	37.85	52.0	1467.0	190.0	496.0	177.0	7.2574	352100.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0	5.6431	341300.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0	3.8462	342200.0

```
#normalization(scaling)
```

```
from sklearn.preprocessing import MinMaxScaler
```

```
scaler=MinMaxScaler()
```

```
df1=pd.DataFrame(scaler.fit_transform(df),columns=df.columns)
```

```
df1
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
0	0.211155	0.567481	0.784314	0.022331	0.019711	0.011168	0.020395	0.539668	0.9051
1	0.212151	0.565356	0.392157	0.180503	0.171349	0.083955	0.186842	0.538027	0.7169
2	0.210159	0.564293	1.000000	0.037260	0.029179	0.017260	0.028783	0.466028	0.7041
3	0.209163	0.564293	1.000000	0.032352	0.036163	0.019431	0.035691	0.354699	0.6825
4	0.209163	0.564293	1.000000	0.041330	0.043148	0.019676	0.042270	0.230776	0.6843
...
15205	0.725100	0.049947	0.078431	0.149245	0.151327	0.067010	0.138158	0.324127	0.588
15206	0.726096	0.049947	0.156863	0.062770	0.064411	0.031544	0.069901	0.270479	1
15207	0.725100	0.048884	0.058824	0.158706	0.194940	0.077303	0.183224	0.258624	0.418
15208	0.725100	0.048884	0.058824	0.232743	0.289306	0.132059	0.278947	0.307244	3
15209	0.724104	0.049947	0.078431	0.143878	0.144187	0.089696	0.148026	0.392753	0.563
...
15210	0.482	1

```
df1.fillna(999,inplace=True)
```

```
y=df1["median_house_value"]
```

```
x=df1.drop("median_house_value",axis=1)
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=42)
```

```
linear_model=LinearRegression()
```

```
lasso_model=Lasso(alpha=0.1)
```

```
ridge_model=Ridge(alpha=0.1)
```

```
linear_model.fit(x_train,y_train)
```

```
LinearRegression()
```

```
lasso_model.fit(x_train,y_train)
```



```

Lasso
Lasso(alpha=0.1)

```

```
ridge_model.fit(x_train,y_train)
```

```

Ridge
Ridge(alpha=0.1)

```

```

print(linear_model.coef_)
print(lasso_model.coef_)
print(ridge_model.coef_)

```

```

[-3.66169442e-01 -3.91829113e-01  8.11385537e-02  6.36677925e-02
-9.63570893e-06 -2.76752624e+00  1.75536617e+00  1.08679972e+00
 2.03408696e-01]
[ 0.00000000e+00 -0.00000000e+00  0.00000000e+00  0.00000000e+00
-1.36304348e-05 -0.00000000e+00  0.00000000e+00  0.00000000e+00
 0.00000000e+00] [-3.65123858e-01 -3.89548815e-01  8.13405031e-02
 7.73727816e-02 -9.92818475e-06 -2.63029307e+00  1.66464762e+00
 1.08561983e+00
 2.04808328e-01]

```

```

print(linear_model.intercept_)
print(lasso_model.intercept_)
print(ridge_model.intercept_)

```

```

0.34227829647618624
0.39767141229141667
0.34050258078362833

```

```

linear_model.score(x_test,y_test)
lasso_model.score(x_test,y_test)
ridge_model.score(x_test,y_test)

```

```

-184.32628077070459

```

```

linear_train_mse=mean_squared_error(y_train,linear_model.predict(x_train))
linear_test_mse=mean_squared_error(y_test,linear_model.predict(x_test))
print(f"Linear Model Training MSE: {linear_train_mse}")
print(f"Linear Model Testing MSE: {linear_test_mse}")

```

```

Linear Model Training MSE: 0.018084107130981655
Linear Model Testing MSE: 9.113051321793758

```

LAB-07

```

#QUESTION-7)Prepare a decision tree model for Iris Dataset using Gini index as criteria
#prepare a decision tree model using the gini index as criteria on the Iris dataset
from sklearn import datasets
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.metrics import accuracy_score
import pandas as pd
import matplotlib.pyplot as plt

df=pd.read_csv("Iris.csv")
df

```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa
...
145	146	6.7	3.0	5.2	2.3	Iris-virginica
146	147	6.3	2.5	5.0	1.9	Iris-virginica
147	148	6.5	3.0	5.2	2.0	Iris-virginica
148	149	6.2	3.4	5.4	2.3	Iris-virginica
149	150	5.9	3.0	5.1	1.8	Iris-virginica

150 rows × 6 columns

```
x=df.drop(['Species','Id'],axis=1)
y=df["Species"]
x
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2
3	4.6	3.1	1.5	0.2
4	5.0	3.6	1.4	0.2
...
145	6.7	3.0	5.2	2.3
146	6.3	2.5	5.0	1.9
147	6.5	3.0	5.2	2.0
148	6.2	3.4	5.4	2.3
149	5.9	3.0	5.1	1.8

150 rows × 4 columns

```
model = DecisionTreeClassifier(criterion='gini')
model
```

DecisionTreeClassifier	?
DecisionTreeClassifier()	

```
gini_impurities={}
import numpy as np
arr=np.array([1,2,3,4,5,6])
print("original array shape: ",arr.shape)
reshaped_arr=arr.reshape(-1,1)
print("Reshaped array shape: ",reshaped_arr.shape)
print(reshaped_arr)
```

original array shape: (6,)
Reshaped array shape: (6, 1)
[[1]
[2]
[3] [4] [5] [6]]

```

for i in range(x.shape[1]):
#fit the classifier with only the current feature
model=model.fit(x.iloc[:,i].values.reshape(-1,1),y)
prob=model.predict_proba(x.iloc[:,i].values.reshape(-1,1))
gini_impurities[i]=1-(prob[:,0]**2+prob[:,1]**2 + prob[:,2]**2).su

best_feature = min(gini_impurities, key=gini_impurities.get)
print(f"Best feature: {best_feature}")

```

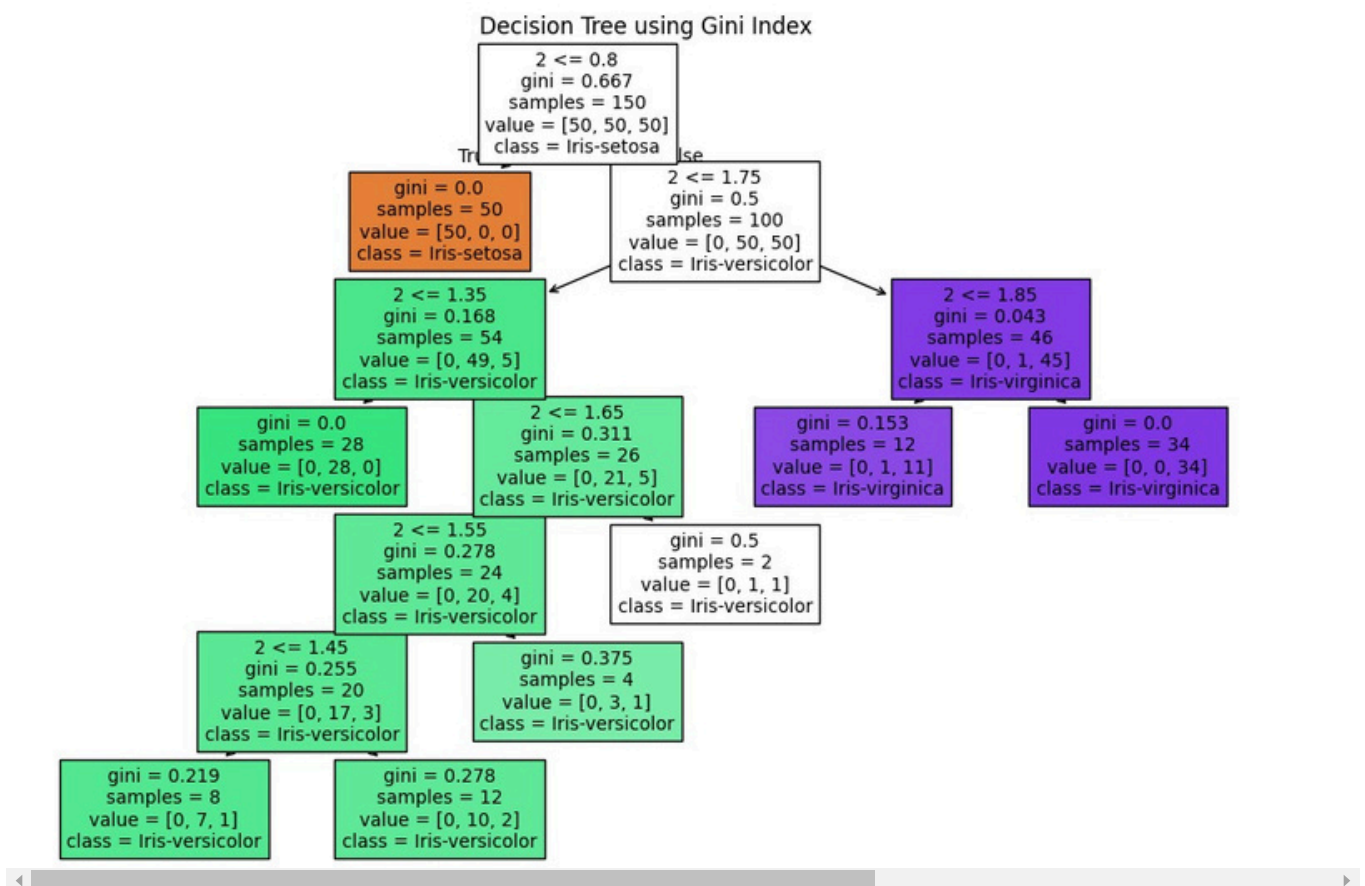
➤ Best feature: 2

```

plt.figure(figsize=(12, 8))
plot_tree(model, filled=True, feature_names=[best_feature], class_na
plt.title("Decision Tree using Gini Index")
plt.show()

```

➤




LAB-08

```

#question-8)Prepare a decision tree model for Iris Dataset using ent
import numpy as np
import pandas as pd
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, plot_tree
import matplotlib.pyplot as plt

```

```
df =pd.read_csv("Iris.csv")
df
```



	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa
...
145	146	6.7	3.0	5.2	2.3	Iris-virginica
146	147	6.3	2.5	5.0	1.9	Iris-virginica
147	148	6.5	3.0	5.2	2.0	Iris-virginica
148	149	6.2	3.4	5.4	2.3	Iris-virginica
149	150	5.9	3.0	5.1	1.8	Iris-virginica

150 rows × 6 columns

```
y = df['Species']
```


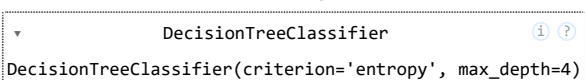
```
x_train , x_test , y_train , y_test = train_test_split(x,y,test_size
```

```
from sklearn import tree
```

```
model = tree.DecisionTreeClassifier(criterion = 'entropy',max_depth
```


```
#fit the tree to the iris dataset
```

```
model.fit(x_train,y_train)
```

```
y_pred = model.predict(x_test)
```

```
print("Accuracy: " , accuracy_score(y_test,y_pred)*100)
```

 Accuracy: 95.55555555555556

```
def plot_decision_tree(model, features_names,class_names):
```

```
plt.figure(figsize=(15,10))
```

```
plot_tree(model,feature_names=features_names,class_names=class_names)
```


```
plt.show()
```

```
model=tree.DecisionTreeClassifier(criterion='entropy',max_depth=4) #
```

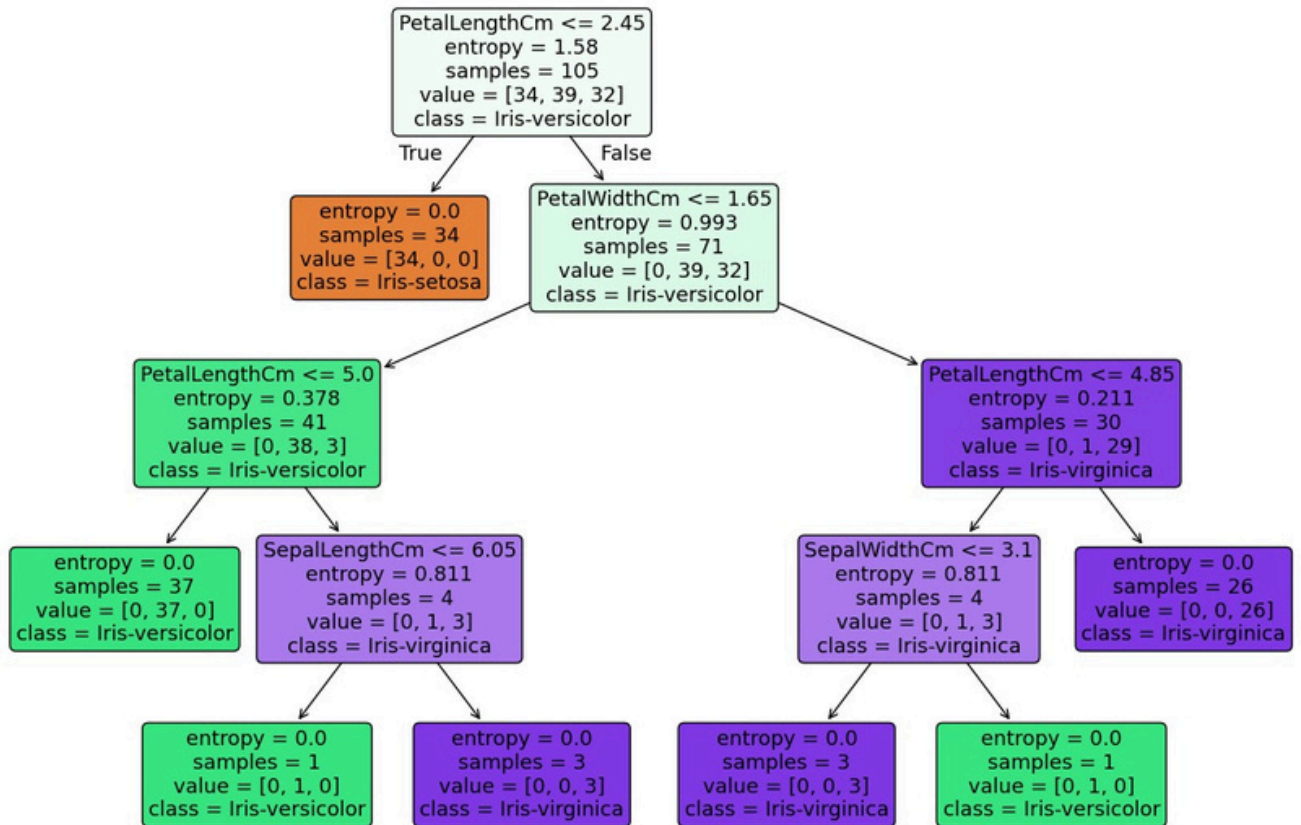
```
model.fit(x_train,y_train)
```

```
y_pred=model.predict(x_test)
```

```
print(accuracy_score(y_test,y_pred)*100)
```

 95.55555555555556

```
plot_decision_tree(model,x.columns,y.unique())
```



LAB-09

#question-9)Prepare a naïve bayes classification model for predictio

```
import pandas as pd
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
import seaborn as sns
```

```
from matplotlib.colors import ListedColormap
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.naive_bayes import GaussianNB
```

```
from sklearn.metrics import confusion_matrix, accuracy_score
```

```
from sklearn.preprocessing import LabelEncoder
```

```
from sklearn import metrics
```

```
from sklearn.metrics import classification_report
```

```
from sklearn.preprocessing import LabelEncoder
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.naive_bayes import GaussianNB
```

```
from sklearn import metrics
```

```
from sklearn.metrics import accuracy_score
```

```

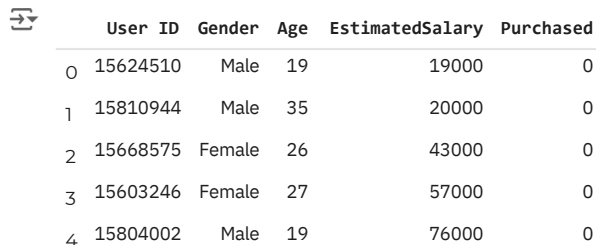
from sklearn.metrics import classification_report
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import confusion_matrix
from sklearn.metrics import f1_score

```

```

df=pd.read_csv("User_Data.csv")
df.head()

```



	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```

df.drop(columns=['User ID'],axis=1,inplace=True)#inplace makes the c
#label encoder is the class which is use to convert a categorical va
#since a ml model is the mathematical model so it understands numeri
le=LabelEncoder()
df['Gender']=le.fit_transform(df['Gender'])

```

```

#split data into dependent and independent variables

```

```

x=df.iloc[:, :-1].values
y=df.iloc[:, -1].values

```

```

X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.25,ra

```

```

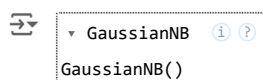
sc=StandardScaler()
X_train=sc.fit_transform(X_train)
X_test=sc.transform(X_test)

```

```

classifier=GaussianNB()
classifier.fit(X_train,y_train)

```



```

GaussianNB()

```

```

#prediction
y_pred=classifier.predict(X_test)
#accuracy
accuracy_score(y_test,y_pred)

```



```

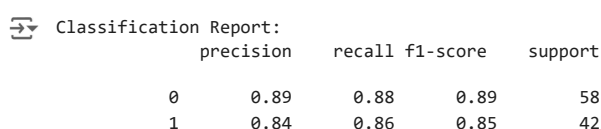
0.87

```

```

print(f'Classification Report:\n{classification_report(y_test,y_pred

```



```

Classification Report:

```

	precision	recall	f1-score	support
0	0.89	0.88	0.89	58
1	0.84	0.86	0.85	42

```
accuracy 0.87 100 macro avg 0.87 0.87 0.87 100
weighted avg 0.87 0.87 0.87 100
```

```
#confusion matrix
cf_matrix=confusion_matrix(y_test,y_pred)
print(cf_matrix)
```

```
[[51 7]
 [ 6 36]]
```

LAB-10

```
#question-10)prepare a naiva bayes model for email classification in
import pandas as pd from sklearn.model_selection import
train_test_split from sklearn.naive_bayes import
MultinomialNB,GaussianNB from sklearn.feature_extraction.text import
CountVectorizer from sklearn.metrics import accuracy_score,f1_score
import matplotlib.pyplot as plt from wordcloud import WordCloud
```

```
df=pd.read_csv("spam.csv",encoding="latin-1")
df.head()
```

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3	ham	U dun say so early hor... U c already then say...	NaN	NaN	NaN
4	ham	Nah I don't think he goes to usf, he lives aro...	NaN	NaN	NaN

```
df=df[['v1','v2']]
df.head()
```

	v1	v2
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

```
df=df.rename(columns={
    'v1':'label',
    'v2':'text'
})
df.head()
```

	label	text
0	ham	Go until jurong point, crazy.. Available only ...
1	ham	Ok lar... Joking wif u oni...
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...
3	ham	U dun say so early hor... U c already then say...
4	ham	Nah I don't think he goes to usf, he lives aro...

```
x=df['text']
y=df['label']
```


```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,ran
```

```
#value_count is used to count the number of unique values in a datas
distribution=y.value_counts()
print(distribution)
```

```
label
ham      4825
spam      747
Name: count, dtype: int64
```

```
distribution.plot(kind='pie',autopct='%1.1f%%')
plt.title('Distribution of Spam and Non-Spam Emails')
plt.show()
```

Distribution of Spam and Non-Spam Emails




```
#generate a word cloud for spam messages
spam_text=' '.join(df[df['label']=='spam']['text'])
spam_text
```

```
'Free entry in 2 a wkly comp to win FA Cup final tkts 21st May 2005. Text FA to 87121 to receive entry question(std txt rate)T&C\'s
apply 08452810075over18\'s FreeMsg Hey there darling it\'s been 3 week\'s now and no word back! I\'d like some fun you up for it st
ill? Tb ok! XxX std chgs to send, â€1.50 to rcv WINNER!! As a valued network customer you have been selected to receivea â€900 priz
e reward! To claim call 09061701461. Claim code KL341. Valid 12 hours only. Had your mobile 11 months or more? U R entitled to Upda
te to the latest colour mobiles with camera for Free! Call The Mobile Update Co FREE on 08002986030 SIX chances to win CASH! From 1 00
to 20,000 pounds txt> CSH11 and send to 87575. Cost 150p/day, 6days, 16+ TsandCs apply Reply HL 4 info URGENT! You have won a 1 week
FREE membership in our â€100,000 Prize Jackpot! Txt the word: CLAIM to No: 81010 T&C www.dbuk.net LCCLTD POBOX 4403LDNW1A7RW18
```

```
spam_wordcloud=WordCloud(width=800,height=400,max_words=100,backgrou
print(spam_wordcloud)
plt.figure(figsize=(10,4))
```




```
plt.imshow(spam_wordcloud)
plt.title('Word Cloud for Spam Messages')
plt.axis('off')
plt
```

 <wordcloud.wordcloud.WordCloud object at 0x7bc0830e8d30>
<module 'matplotlib.pyplot' from '/usr/local/lib/python3.10/dist-packages/matplotlib/pyplot.py'>





```
ham_text=' '.join(df[df['label']=='ham']['text'])
ham_text
```

 'Go until jurong point, crazy.. Available only in bugis n great world la e buffet... Cine there got amore wat... Ok lar... Joking w if u oni... U dun say so early hor... U c already then say... Nah I don't think he goes to usf, he lives around here though Even m y brother is not like to speak with me. They treat me like aids patent. As per your request 'Melle Melle (Oru Minnaminunginte Nuru ngu Vettam)' has been set as your callertune for all Callers. Press *9 to copy your friends Callertune I'm gonna be home soon and i don't want to talk about this stuff anymore tonight, k? I've cried enough today. I've been searching for the right words to thank you for this breather. I promise i wont take your help for granted and will fulfil my promise. You have been wonderful and a blessing at all times. I HAVE A DATE ON SUNDAY WITH WILL!! Oh k...i'm watching here:) Eh u remember how 2 spell his name... Yes i di

```
ham_wordcloud=WordCloud(width=800,height=400,max_words=100,backgroun
```

```
plt.subplot(1,2,2)
plt.imshow(ham_wordcloud)
plt.title("Word Cloud for ham Messages")
plt.axis("off")
```

 (-0.5, 799.5, 399.5, -0.5)



```
#count vectorizer is a text processing technique used in natural la
vectorizer=CountVectorizer()
x_train=vectorizer.fit_transform(x_train)
x_test=vectorizer.transform(x_test)
```

```
#alpha is the regularization parametre and prevent overfitting
model_multinomial=MultinomialNB(alpha=0.8,fit_prior=True,force_alpha
model_multinomial.fit(x_train,y_train)
```

↗ MultinomialNB ⓘ ?
MultinomialNB(alpha=0.8)

```
model_guassian=GaussianNB()  
model_guassian.fit(x_train.toarray(),y_train)
```

↗ GaussianNB ⓘ ?
GaussianNB()

```
y_pred_multinomial=model_multinomial.predict(x_test)  
accuracy_multinomial=accuracy_score(y_test,y_pred_multinomial)  
print("Accuracy for multinomial model is:",accuracy_multinomial)
```

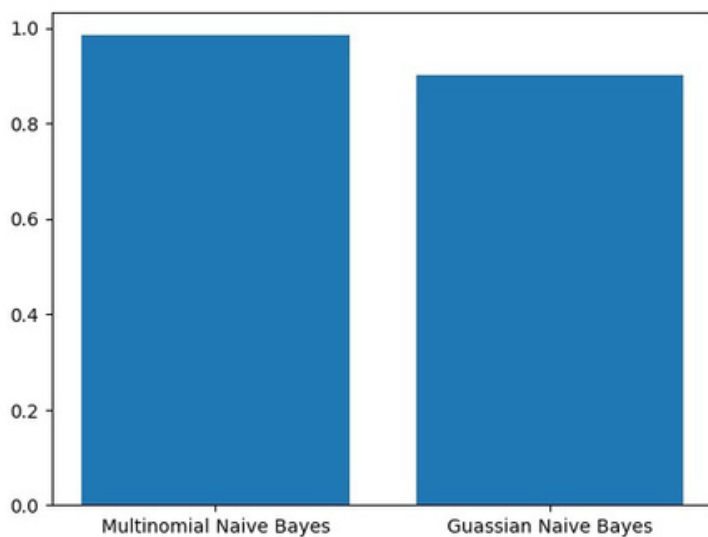
↗ Accuracy for multinomial model is: 0.9838565022421525

```
y_pred_guassian=model_guassian.predict(x_test.toarray())  
accuracy_guassian=accuracy_score(y_test,y_pred_guassian)  
print("Accuracy for guassian model is:",accuracy_guassian)
```

↗ Accuracy for guassian model is: 0.9004484304932735

```
methods=["Multinomial Naive Bayes","Guassian Naive Bayes"]  
scores=[accuracy_multinomial,accuracy_guassian]  
plt.bar(methods,scores)
```

↗ <BarContainer object of 2 artists>



□ LAB-11

```
#Question-11)prepare a model for prediction of prostate cancer using  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.model_selection import train_test_split
```

```
df=pd.read_csv("prostate.csv")
```

```
df
```

	lcavol	lweight	age	lbph	lcp	gleason	pgg45	lpsa	Target
0	-0.579818	2.769459	50	-1.386294	-1.386294	6	0	-0.430783	0
1	-0.994252	3.319626	58	-1.386294	-1.386294	6	0	-0.162519	0
2	-0.510826	2.691243	74	-1.386294	-1.386294	7	20	-0.162519	0
3	-1.203973	3.282789	58	-1.386294	-1.386294	6	0	-0.162519	0
4	0.751416	3.432373	62	-1.386294	-1.386294	6	0	0.371564	0
...									
92	2.830268	3.876396	68	-1.386294	1.321756	7	60	4.385147	1
93	3.821004	3.896909	44	-1.386294	2.169054	7	40	4.684443	1
94	2.907447	3.396185	52	-1.386294	2.463853	7	10	5.143124	1
95	2.882564	3.773910	68	1.558145	1.558145	7	80	5.477509	1
96	3.471966	3.974998	68	0.438255	2.904165	7	20	5.582932	1

97 rows × 9 columns

```
df.shape
```

```
(97, 9)
```

```
x=df.drop("Target",axis=1)
```

```
y=df["Target"]
```

```
#feature scaling(to convert the values between 0 and 1)
```

```
scaler=StandardScaler()
```

```
df1=pd.DataFrame(scaler.fit_transform(x),columns=df.columns[:-1])
```

```
df1.head()
```

	lcavol	lweight	age	lbph	lcp	gleason	pgg45	lpsa
0	-1.645861	-2.016634	-1.872101	-1.030029	-0.867655	-1.047571	-0.868957	-2.533318
1	-1.999313	-0.725759	-0.791989	-1.030029	-0.867655	-1.047571	-0.868957	-2.299712
2	-1.587021	-2.200154	1.368234	-1.030029	-0.867655	0.344407	-0.156155	-2.299712
3	-2.178174	-0.812191	-0.791989	-1.030029	-0.867655	-1.047571	-0.868957	-2.299712
4	-0.510513	-0.461218	-0.251933	-1.030029	-0.867655	-1.047571	-0.868957	-1.834631

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,ran
```

```
knn_model=KNeighborsClassifier(n_neighbors=1)
```

```
knn_model.fit(x_train,y_train)
```

```
KNeighborsClassifier
```

```
KNeighborsClassifier(n_neighbors=1)
```

```
y_pred=knn_model.predict(x_test)
```

```
print(confusion_matrix(y_test,y_pred))
```

```
[[18 4]
 [ 6 2]]
```

```
print(classification_report(y_test,y_pred))
```

```
precision recall f1-score support
```

```
0 0.75 0.82 0.78 22
```

```
1 0.33 0.25 0.29 8
```

```
accuracy 0.67 30
```

```
macro avg 0.54 0.53 0.53 30
```


```
weighted avg 0.64 0.67 0.65 30
```

```
#elbow method for calculating k
error_rate=[]
```

```
for i in range(1,40):
knn=KNeighborsClassifier(n_neighbors=i)
knn.fit(x_train,y_train)
new_y_pred=knn.predict(x_test)
error_rate.append(np.mean(new_y_pred!=y_test))
```

```
plt.figure(figsize=(16,12))
plt.plot(error_rate,color='blue',linestyle='dashed',marker='o',marke
```

```
plt.title('Error Rate vs K Value')
plt.xlabel('K')
plt.ylabel('Error Rate')
plt.show()
```




□ LAB-12

#Question-12)prepare a model for prediction of survival from Titanic


```
import pandas as pd
import numpy as np
import warnings
warnings.filterwarnings('ignore')
import matplotlib.pyplot as plt
from sklearn.preprocessing import LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, class
```

```
df=pd.read_csv("Titanic-Dataset.csv")
df.head()
```




	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S

```
df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    bool
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age          714 non-null    float64
6   Parch       891 non-null    int64
7   Ticket       891 non-null    object
8   Fare         891 non-null    float64
9   Cabin        204 non-null    object
10  Embarked     891 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
df.dropna(subset=['Survived'])
```



	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S
...
886	887	0	2	Montvila, Rev. Juozas	male	27.0	0	0	211536	13.0000	NaN	S
887	888	1	1	Graham, Miss. Margaret Edith	female	19.0	0	0	112053	30.0000	B42	S
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"	female	NaN	1	2	W./C. 6607	23.4500	NaN	S
889	890	1	1	Behr, Mr. Karl Howell	male	26.0	0	0	111369	30.0000	C148	C


```
df.shape
```



```
(891, 12)
```

```
x=df[['Pclass','Sex','Age','SibSp','Parch','Fare']]
y=df['Survived']
```

```
x.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 6 columns):
#   Column Non-Null Count  Dtype
---  -
0   Pclass 891 non-null      int64
```

```

1 Sex      891 non-null object
2 Age      714 non-null float64
3 SibSp    891 non-null int64
4 Parch    891 non-null int64
5 Fare     891 non-null float64

dtypes: float64(2), int64(3), object(1)
memory usage: 41.9+ KB


```

```

le=LabelEncoder()
x['Sex']=le.fit_transform(x['Sex'])

```


```
x.head()
```



	Pclass	Sex	Age	SibSp	Parch	Fare
0	3	1	22.0	1	0	7.2500
1	1	0	38.0	1	0	71.2833
2	3	0	26.0	0	0	7.9250
3	1	0	35.0	1	0	53.1000
4	3	1	35.0	0	0	8.0500

```
x['Age']=x['Age'].fillna(x['Age'].mean())
```

```
x['Age']
```



	Age
0	22.000000
1	38.000000
2	26.000000
3	35.000000
4	35.000000
...	...
886	27.000000
887	19.000000
888	29.699118
889	26.000000
890	32.000000


891 rows × 1 columns

dtype: float64

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,ran
```

```
model=RandomForestClassifier(n_estimators=100,random_state=42)
```

```
model.fit(x_train,y_train)
```



▼	RandomForestClassifier	i	?
RandomForestClassifier(random_state=42)			

```
y_pred=model.predict(x_test)
```